

# Yolah Board Game

Building a Two-Player Perfect-Information Game with AI  
Players

Pascal Garcia

December 9, 2025

MrCoder

C'est en forgeant qu'on devient  
forgeron

*À Sarah, Hugo et Célia* ❤️

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	The Yolah Game . . . . .	1
1.1.1	Game Rules . . . . .	1
1.1.2	Interesting Characteristics of Yolah for Developing AIs . . . .	5
1.1.3	Game Example . . . . .	5
1.1.4	What's Next . . . . .	6
<b>2</b>	<b>Game Engine</b>	<b>11</b>
2.1	Data Structures . . . . .	11
2.2	Game Over Test . . . . .	15
2.3	Generating Possible Piece Moves . . . . .	19
2.3.1	Magic Perfect Hashing Function . . . . .	19
2.3.2	Magic Bitboards for Piece Moves . . . . .	28
2.3.3	Making and Unmaking Moves . . . . .	28
2.4	Complete Commented Game Board Code . . . . .	28
2.5	Testing with Random Games . . . . .	33
2.6	What's Next . . . . .	33
<b>3</b>	<b>AI Players</b>	<b>35</b>
<b>4</b>	<b>Monte Carlo Player</b>	<b>37</b>
<b>5</b>	<b>MCTS Player</b>	<b>39</b>
<b>6</b>	<b>Minmax Player</b>	<b>41</b>
<b>7</b>	<b>Minmax with Neural Network Player</b>	<b>43</b>
<b>8</b>	<b>AI Tournament</b>	<b>45</b>
<b>9</b>	<b>Conclusion</b>	<b>47</b>
	<b>Acronymes</b>	<b>49</b>
	<b>Bibliography</b>	<b>51</b>



# List of Figures

1.1	The Pingouins game box . . . . .	1
1.2	The initial configuration of the Yolah game . . . . .	2
1.3	Possible moves (small black crosses) for the black piece located on square d5 . . . . .	3
1.4	Black just moved from d5 to b7. The starting square d5 becomes inaccessible and impassable for the rest of the game . . . . .	3
1.5	Possible moves (small white crosses) for the white piece located on square e5 . . . . .	4
1.6	White just moved from e5 to f5. The starting square e5 becomes inaccessible and impassable for the rest of the game. The score is one point each (each player has moved once) . . . . .	4
1.7	Game example between two AIs - moves 1 to 19 . . . . .	7
1.8	Game example between two AIs - moves 20 to 39 . . . . .	8
1.9	Game example between two AIs - moves 40 to 56. White wins 32 to 24 . . . . .	9
2.1	Board configuration corresponding to move 21 of the game given as an example in the previous chapter (Figure 1.8b) . . . . .	12





# List of Tables

2.1	Position of each board square in the bitboard . . . . .	11
2.2	Black piece positions . . . . .	13
2.3	White piece positions . . . . .	13
2.4	Destroyed square positions (holes) . . . . .	13
2.5	Black and white piece positions obtained by computing: <code>black   white</code>	15
2.6	Position of each board square in the bitboard. Note that for a square not in rank 8, <code>bit<sub>i+8</sub></code> corresponds to the square north of <code>bit<sub>i</sub></code> . . . . .	16
2.7	Game board before applying <code>shift&lt;NORTH&gt;</code> . . . . .	17
2.8	Game board after applying <code>shift&lt;NORTH&gt;</code> . . . . .	17
2.9	Game board . . . . .	18
2.10	Positions around the pieces . . . . .	18



# Chapter 1

## Introduction

### 1.1 The Yolah Game

I created the Yolah game to illustrate effective techniques for implementing board games and artificial intelligences for my students. I was inspired by the Pingouins game, whose box you can see in Figure 1.1 (I highly recommend it ☺)




Figure 1.1: The Pingouins game box

#### Important

I have done my best with my current knowledge (*ars longa, vita brevis*) to implement my game and the associated AIs. But like any good scientist, you should look at my work with a critical eye. I wrote the book in French (easier for me) and asked an AI assistant (Claude [1]) to translate it for me.

I will now describe the rules of the game, then I will explain why I chose these rules, I will give an example of a game between two AIs and then I will present the rest of the book.

#### 1.1.1 Game Rules

The Yolah game board is shown in Figure 1.2. You can see four black pieces and four white pieces placed symmetrically. Black starts by choosing one of their four




Figure 1.2: The initial configuration of the Yolah game

pieces. A piece can never disappear from the board because Yolah is a game without captures. A piece moves in all eight directions as far as it wishes as long as it is not blocked by another piece or a hole (a concept we will soon discuss). For example, if black chooses to move their piece located at d5, the squares where it can land are indicated by small black crosses in Figure 1.3.

Now, if the black piece at d5 moves to b7, which we will denote as d5:b7, we get the configuration shown in Figure 1.4. Notice that the starting square of the black piece disappears and becomes a hole! This square (this hole) becomes inaccessible and impassable for the rest of the game! This will create opportunities to block the opponent and try to create areas where the opponent cannot go.

A move earns one point for the player who just moved. For example, in the configuration of Figure 1.4, the black player has one point and the white player who has not yet moved has zero points. The goal of the game is quite simple to summarize: you must move longer than your opponent!

Now it is white's turn to play. They must decide which white piece they will move. Suppose it is the piece at e5. The possible moves for this white piece are shown in Figure 1.5. If white decides to make the move e5:f5, we end up in the configuration of Figure 1.6 and the score is one point each (each player has played one move).

To summarize, the rules of Yolah are as follows:

- The game is a two-player game (black and white) played in turns.




Figure 1.3: Possible moves (small black crosses) for the black piece located on square d5




Figure 1.4: Black just moved from d5 to b7. The starting square d5 becomes inaccessible and impassable for the rest of the game




Figure 1.5: Possible moves (small white crosses) for the white piece located on square e5




Figure 1.6: White just moved from e5 to f5. The starting square e5 becomes inaccessible and impassable for the rest of the game. The score is one point each (each player has moved once)

- Each player has four pieces.
- On their turn, the player chooses one of the pieces that can still move; if no piece can move, they pass their turn (we will denote by the move `a1:a1` to skip one's turn).
- They must move the chosen piece in one of the eight directions, as many squares as desired, but must not land on or be blocked by a piece or a hole.
- After moving the chosen piece, the starting square of the move becomes a hole and can no longer be crossed or landed on.
- After each move, the player earns one point.
- The game ends when both players can no longer move.
- The player with the most points wins the game.
- If both players have the same number of points, the game is declared a draw.

### 1.1.2 Interesting Characteristics of Yolah for Developing AIs

I chose the above rules for Yolah, on the one hand because I liked the dynamics of the Penguin game, but also because there are no cycles in the game, so there is no need for special rules to prevent a game from never ending. The number of moves available to each player is quite large at the beginning (but reasonable)<sup>1</sup>, but it will decrease little by little with the appearance of holes<sup>2</sup>. This allows the AIs to look ahead a fairly large number of moves.

I also wanted to be able to reuse concepts used for the efficient implementation of chess, and the size of the board and the movement of the pieces (same way of moving as a Queen in chess) allow me to do that.

### 1.1.3 Game Example

To get an idea of how a Yolah game unfolds, we will have two artificial intelligences play against each other. The first AI will be based on Monte Carlo Tree Search and the second will be based on Minimax with a neural network. We will study both of these AIs later in the book. The second AI is stronger and you will see its zone isolation strategy in action!

The progression of the game is described in Figures 1.7, 1.8 and 1.9.

The white AI estimates that it is winning starting from move 10 (see Figure 1.7k). We can see at move 30 (see Figure 1.8k) that it has successfully isolated a zone where black can no longer access. At move 32 (see Figure 1.8m) it moves one of its pieces out of the isolated zone because the other piece will be able to collect all the points from

---

<sup>1</sup>56 possible moves for the black player at the start of the game.

<sup>2</sup>The number of possible moves does not necessarily decrease after each move; there are configurations where the player has more than 56 available moves.

that zone. It is more useful to use the other piece to gather points elsewhere. Note that starting from move 47 (see Figure 1.9h onward), black has no more available moves and must therefore pass their turn.

The game is won by the white player 32 to 24, which is a very good score because the Yolah game seems to favor black.

### 1.1.4 What's Next

In the next chapter, we will study the implementation of the Yolah game in [2]. This implementation is designed to be efficient because it will be important for the AIs to be able to play many games per second; their level of play will depend on it.

Chapter 3 describes the common interface for our different AIs. Chapter 4 presents a very simple AI based on Monte Carlo search. The next AI, described in Chapter ??, is an evolution of the previous one and will allow, unlike the Monte Carlo AI, the development of a game tree. Note that these two AIs will not require heuristics provided by humans and only need the rules of the game. Chapter 6 presents an AI based on minimax tree search with heuristics provided by humans. The heuristic used by the AI is a linear combination of the heuristics provided by humans; the weights of each heuristic in this linear combination are learned using a genetic algorithm. Our last and strongest AI is presented in Chapter 7. A neural network is used instead of heuristics. This neural network is trained on a set of games played by the previous AI.

Chapter 8 then evaluates all these AIs by having them compete in a tournament. We will then conclude and propose different directions for creating other artificial players.

Happy reading!






Figure 1.7: Game example between two AIs - moves 1 to 19




Figure 1.8: Game example between two AIs - moves 20 to 39




Figure 1.9: Game example between two AIs - moves 40 to 56. White wins 32 to 24



# Chapter 2

## Game Engine

We will implement game management in C++, striving to create an efficient implementation capable of running as many games per second as possible—this will be crucial for developing high-level AIs. We will test our implementation by generating as many random games as possible within a given time at the end of this chapter. We will draw inspiration from the excellent Stockfish chess engine [3] for Yolah’s data structures.

### 2.1 Data Structures

We will represent the positions of black pieces, white pieces, and destroyed squares using 64-bit unsigned integers: `uint64_t`. We call these integers bitboards. We use `uint64_t` because the Yolah board contains 64 squares, giving us one bit per square. Table 2.1 shows the position of each board square in the bitboard. This information is represented in code by the enumeration in Listing 1<sup>1</sup>.

Table 2.1: Position of each board square in the bitboard

8	bit <sub>56</sub>	bit <sub>57</sub>	bit <sub>58</sub>	bit <sub>59</sub>	bit <sub>60</sub>	bit <sub>61</sub>	bit <sub>62</sub>	bit <sub>63</sub>
7	bit <sub>48</sub>	bit <sub>49</sub>	bit <sub>50</sub>	bit <sub>51</sub>	bit <sub>52</sub>	bit <sub>53</sub>	bit <sub>54</sub>	bit <sub>55</sub>
6	bit <sub>40</sub>	bit <sub>41</sub>	bit <sub>42</sub>	bit <sub>43</sub>	bit <sub>44</sub>	bit <sub>45</sub>	bit <sub>46</sub>	bit <sub>47</sub>
5	bit <sub>32</sub>	bit <sub>33</sub>	bit <sub>34</sub>	bit <sub>35</sub>	bit <sub>36</sub>	bit <sub>37</sub>	bit <sub>38</sub>	bit <sub>39</sub>
4	bit <sub>24</sub>	bit <sub>25</sub>	bit <sub>26</sub>	bit <sub>27</sub>	bit <sub>28</sub>	bit <sub>29</sub>	bit <sub>30</sub>	bit <sub>31</sub>
3	bit <sub>16</sub>	bit <sub>17</sub>	bit <sub>18</sub>	bit <sub>19</sub>	bit <sub>20</sub>	bit <sub>21</sub>	bit <sub>22</sub>	bit <sub>23</sub>
2	bit <sub>8</sub>	bit <sub>9</sub>	bit <sub>10</sub>	bit <sub>11</sub>	bit <sub>12</sub>	bit <sub>13</sub>	bit <sub>14</sub>	bit <sub>15</sub>
1	bit <sub>0</sub>	bit <sub>1</sub>	bit <sub>2</sub>	bit <sub>3</sub>	bit <sub>4</sub>	bit <sub>5</sub>	bit <sub>6</sub>	bit <sub>7</sub>
	a	b	c	d	e	f	g	h

---

```
1 enum Square : int8_t {  
2     SQ_A1, SQ_B1, SQ_C1, SQ_D1, SQ_E1, SQ_F1, SQ_G1, SQ_H1,  
3     SQ_A2, SQ_B2, SQ_C2, SQ_D2, SQ_E2, SQ_F2, SQ_G2, SQ_H2,
```

---

<sup>1</sup>By default, the enumeration starts at value 0, so `SQ_A1` equals 0, `SQ_B1` equals 1, ...




Table 2.2: Black piece positions

8	.	•	.	.	.	.	.	.
7	.	.	.	.	.	.	.	•
6	•	.	.	.	.	.	.	.
5	.	.	.	.	.	.	.	.
4	.	.	.	.	.	•	.	.
3	.	.	.	.	.	.	.	.
2	.	.	.	.	.	.	.	.
1	.	.	.	.	.	.	.	.
	a	b	c	d	e	f	g	h

Table 2.3: White piece positions

8	.	.	.	.	.	.	.	.
7	.	.	.	.	.	○	.	.
6	.	.	○	.	.	.	.	.
5	.	.	.	.	.	.	.	.
4	.	.	.	.	.	.	.	.
3	.	.	.	.	.	.	.	.
2	.	.	.	○	.	○	.	.
1	.	.	.	.	.	.	.	.
	a	b	c	d	e	f	g	h

Table 2.4: Destroyed square positions (holes)

8	×	.	.	.	×	.	.	×
7	.	.	.	.	.	.	.	.
6	.	.	.	×	×	×	.	.
5	.	.	.	×	×	×	×	×
4	.	.	.	×	×	.	×	.
3	×	.	.	.	.	×	.	.
2	×	×	.	.	.	.	.	.
1	×	.	.	.	.	.	×	×
	a	b	c	d	e	f	g	h

Bitboards allow us to efficiently manipulate the board using bitwise operations while requiring minimal memory to represent it.

We represent the game board with the `Yolah` class, an excerpt of which is given in Listing 2.

---

```
1 constexpr uint64_t BLACK_INITIAL_POSITION =
2 0b10000000'00000000'00000000'00001000'00010000'00000000'00000000'00000001;
3 constexpr uint64_t WHITE_INITIAL_POSITION =
4 0b00000001'00000000'00000000'00010000'00001000'00000000'00000000'10000000;
5 class Yolah {
6     uint64_t black = BLACK_INITIAL_POSITION;
7     uint64_t white = WHITE_INITIAL_POSITION;
8     uint64_t holes = 0;
9     uint8_t black_score = 0;
10    uint8_t white_score = 0;
11    uint8_t ply = 0;
12 public:
13     // ...
14 };
```

---

Listing 2: Attributes of the `Yolah` class representing the game board

The attributes of the `Yolah` class are:

- Line 6, `black`: the bitboard for black pieces.
- Line 7, `white`: the bitboard for white pieces.
- Line 8, `holes`: the bitboard for holes (destroyed squares).
- Line 9, `black_score`: the score, or number of moves, of the black player.
- Line 10, `white_score`: the score, or number of moves, of the white player.
- Line 11, `ply`: the number of moves played by both players since the start of the game.

A `Yolah` object will occupy `sizeof(Yolah) == 32` bytes in memory. Note that due to padding<sup>2</sup>, it would not have been wise to write, for example,

---

```
1 class Yolah {
2     uint64_t black = BLACK_INITIAL_POSITION;
```

---

<sup>2</sup>[https://en.wikipedia.org/wiki/Data\\_structure\\_alignment](https://en.wikipedia.org/wiki/Data_structure_alignment)



---

```

3     uint8_t black_score = 0;
4     uint64_t white = WHITE_INITIAL_POSITION;
5     uint8_t white_score = 0;
6     uint64_t holes = 0;
7     uint8_t ply = 0;
8 public:
9     // ...
10 };

```

---

because with this implementation, we would have `sizeof(Yolah) == 48` bytes.

The bitboard representation allows for efficiently obtaining game information. For example, to see the squares occupied by black and white pieces, simply perform: `black | white`. For the positions in Tables 2.2 and 2.3, we would obtain in a single highly efficient operation<sup>3</sup> the black and white piece positions shown in Table 2.5. In binary notation, we get:

```

black | white
== 00000010100000000000000100000000001000000000000000000000000000 |
   00000000001000000000010000000000000000000000000000001010000000000
== 0000001010100000000001010000000000100000000000001010000000000

```

Table 2.5: Black and white piece positions obtained by computing: `black | white`

8	.	•	.	.	.	.	.	.
7	.	.	.	.	.	○	.	•
6	•	.	○	.	.	.	.	.
5	.	.	.	.	.	.	.	.
4	.	.	.	.	.	•	.	.
3	.	.	.	.	.	.	.	.
2	.	.	.	○	.	○	.	.
1	.	.	.	.	.	.	.	.
	a	b	c	d	e	f	g	h

We will continue using bitwise operations throughout this chapter, particularly for testing game over conditions and generating possible moves.

## 2.2 Game Over Test

To test for game over, we will need some enumerations and constants (Listings 3 and 5) and the `shift` function (Listing 4). The `NORTH` constant in the `Direction` enumeration (Listing 3) equals 8. Why this value? Let us revisit below, see Table 2.6, the table

---

<sup>3</sup>You can find information on instruction latency and throughput for various processors at: [https://agner.org/optimize/instruction\\_tables.pdf](https://agner.org/optimize/instruction_tables.pdf)

from the previous chapter showing the position of each bit in a bitboard on the board. We can see that adding 8 to the bit number in any square gives us the bit number of the square to the north (unless we exit the board). The same applies to the other constant values in the `Direction` enumeration (Listing 3).

Table 2.6: Position of each board square in the bitboard. Note that for a square not in rank 8,  $\text{bit}_{i+8}$  corresponds to the square north of  $\text{bit}_i$

8	bit <sub>56</sub>	bit <sub>57</sub>	bit <sub>58</sub>	bit <sub>59</sub>	bit <sub>60</sub>	bit <sub>61</sub>	bit <sub>62</sub>	bit <sub>63</sub>
7	bit <sub>48</sub>	bit <sub>49</sub>	bit <sub>50</sub>	bit <sub>51</sub>	bit <sub>52</sub>	bit <sub>53</sub>	bit <sub>54</sub>	bit <sub>55</sub>
6	bit <sub>40</sub>	bit <sub>41</sub>	bit <sub>42</sub>	bit <sub>43</sub>	bit <sub>44</sub>	bit <sub>45</sub>	bit <sub>46</sub>	bit <sub>47</sub>
5	bit <sub>32</sub>	bit <sub>33</sub>	bit <sub>34</sub>	bit <sub>35</sub>	bit <sub>36</sub>	bit <sub>37</sub>	bit <sub>38</sub>	bit <sub>39</sub>
4	bit <sub>24</sub>	bit <sub>25</sub>	bit <sub>26</sub>	bit <sub>27</sub>	bit <sub>28</sub>	bit <sub>29</sub>	bit <sub>30</sub>	bit <sub>31</sub>
3	bit <sub>16</sub>	bit <sub>17</sub>	bit <sub>18</sub>	bit <sub>19</sub>	bit <sub>20</sub>	bit <sub>21</sub>	bit <sub>22</sub>	bit <sub>23</sub>
2	bit <sub>8</sub>	bit <sub>9</sub>	bit <sub>10</sub>	bit <sub>11</sub>	bit <sub>12</sub>	bit <sub>13</sub>	bit <sub>14</sub>	bit <sub>15</sub>
1	bit <sub>0</sub>	bit <sub>1</sub>	bit <sub>2</sub>	bit <sub>3</sub>	bit <sub>4</sub>	bit <sub>5</sub>	bit <sub>6</sub>	bit <sub>7</sub>
	a	b	c	d	e	f	g	h

---

```

1  enum Direction : int8_t {
2      NORTH = 8,
3      EAST  = 1,
4      SOUTH = -NORTH,
5      WEST  = -EAST,
6      NORTH_EAST = NORTH + EAST,
7      SOUTH_EAST = SOUTH + EAST,
8      SOUTH_WEST = SOUTH + WEST,
9      NORTH_WEST = NORTH + WEST
10 };

```

---

Listing 3: Directions

---

```

1  template<Direction D>
2  constexpr uint64_t shift(uint64_t b) {
3      if constexpr (D == NORTH)
4          return b << NORTH;
5      else if constexpr (D == SOUTH)
6          return b >> -SOUTH;
7      else if constexpr (D == EAST)
8          return (b & ~FileHBB) << EAST;
9      else if constexpr (D == WEST)
10         return (b & ~FileABB) >> -WEST;

```

---

```

11     else if constexpr (D == NORTH_EAST)
12         return (b & ~FileHBB) << NORTH_EAST;
13     else if constexpr (D == NORTH_WEST)
14         return (b & ~FileABB) << NORTH_WEST;
15     else if constexpr (D == SOUTH_EAST)
16         return (b & ~FileHBB) >> -SOUTH_EAST;
17     else if constexpr (D == SOUTH_WEST)
18         return (b & ~FileABB) >> -SOUTH_WEST;
19     else return 0;
20 }

```

---

Listing 4: Shift by direction

The `shift(uint64_t b)` function given in Listing 4 shifts all 1-bits in the bitboard parameter in direction `D`<sup>4</sup>. Given

`black == 1000000000000000000000000100000100000000000000000000000000000001`

the bitboard shown in Table 2.7. To shift the black pieces one square north, we perform: `shift<NORTH>(black)`. We then get the bitboard

`1000000000000000000000000000010000010000000000000000000000000001`

shown in Table 2.8. Note that the piece on h8 is no longer on the board.

Table 2.7: Game board before applying `shift<NORTH>`

8	.	.	.	.	.	.	.	●
7	.	.	.	.	.	.	.	.
6	.	.	.	.	.	.	.	.
5	.	.	.	●	.	.	.	.
4	.	.	.	.	.	●	.	.
3	.	.	.	.	.	.	.	.
2	.	.	.	.	.	.	.	.
1	●	.	.	.	.	.	.	.
	a	b	c	d	e	f	g	h

Table 2.8: Game board after applying `shift<NORTH>`

8	.	.	.	.	.	.	.	.
7	.	.	.	.	.	.	.	.
6	.	.	.	●	.	.	.	.
5	.	.	.	.	.	●	.	.
4	.	.	.	.	.	.	.	.
3	.	.	.	.	.	.	.	.
2	●	.	.	.	.	.	.	.
1	.	.	.	.	.	.	.	.
	a	b	c	d	e	f	g	h

We can thus easily obtain all squares directly adjacent to the pieces on a given board by performing the following operation:

---

```

1  shift<NORTH>(board) | shift<SOUTH>(board) | shift<EAST>(board) |
2  shift<WEST>(board) | shift<NORTH_EAST>(board) | shift<NORTH_WEST>(board) |
3  shift<SOUTH_EAST>(board) | shift<SOUTH_WEST>(board)

```

---

<sup>4</sup>Thanks to `if constexpr`, when we call `shift<NORTH>(b)` for example, the compiler will transform the code to: `return b << NORTH;`.

For the board in Table 2.9, we get the positions represented by stars in Table 2.10.

Table 2.9: Game board

8	○	.	.	.	.	.	.	●
7	.	●	.	.	.	.	.	.
6	.	.	.	.	.	.	.	.
5	.	.	.	×	○	.	.	.
4	.	.	.	×	×	●	.	.
3	.	.	.	.	.	.	.	.
2	.	.	.	○	.	.	.	.
1	●	.	.	.	.	.	.	○
	a	b	c	d	e	f	g	h

Table 2.10: Positions around the pieces

8	*	*	*	.	.	.	*	.
7	*	*	*	.	.	.	*	*
6	*	*	*	*	.	*	.	.
5	.	.	.	*	*	*	*	.
4	.	.	.	*	*	*	*	.
3	.	.	*	*	*	*	*	.
2	*	*	*	.	*	.	*	*
1	.	*	*	*	*	.	*	.
	a	b	c	d	e	f	g	h

In the `shift` function code (Listing 4), we can see the use of constants `FileHBB` and `FileABB` (BB for bitboard). These constants allow us to mask certain bits that would end up in incorrect positions after shifting. For example, if we shift the bitboard shown in Table 2.9 in the `EAST` direction, the white pawn on `h1` would end up on `a2`. To avoid this, before shifting, we eliminate elements from column `h` so they don't end up in column `a` (line 8 of Listing 4). The constants `FileABB` and `FileHBB` along with the rank and file enumerations are defined in Listing 5.

---

```

1  enum File : uint8_t {
2      FILE_A, FILE_B, FILE_C, FILE_D, FILE_E, FILE_F, FILE_G, FILE_H,
3      FILE_NB
4  };
5  enum Rank : uint8_t {
6      RANK_1, RANK_2, RANK_3, RANK_4, RANK_5, RANK_6, RANK_7, RANK_8,
7      RANK_NB
8  };
9  constexpr uint64_t FileABB = 0x0101010101010101;
10 // 0x0101010101010101
11 //== 0b000000001000000010000000100000001000000010000000100000001
12 //      a8      a7      a6      a5      a4      a3      a2      a1
13
14 constexpr uint64_t FileHBB = FileABB << 7;
15 // 0x8080808080808080
16 //== 0b10000000100000001000000010000000100000001000000010000000
17 //      h8      h7      h6      h5      h4      h3      h2      h1

```

---

Listing 5: Files and ranks

The function that tests for game over is called `game_over` and is described in Listing

6. Its implementation is straightforward:

- On line 2, we retrieve in `possible` the bitboard with 1s in free positions.
- On line 3, we create the `players` bitboard containing 1s in positions occupied by either player.
- Lines 4 to 8 create the `around_players` bitboard containing 1s at positions around each player's pieces.
- Finally, on line 9, we test whether no free position exists adjacent to either player. The bitwise AND will keep a 1 in a result bit if and only if there is a 1 at that position in both the `possible` and `around_players` bitboards—meaning the corresponding square is both free and accessible by a player. If `around_players & possible` equals 0, it means no position is both adjacent to a player and free.

---

```
1  bool Yolah::game_over() const {
2      uint64_t possible = ~holes & ~black & ~white;
3      uint64_t players  = black | white;
4      uint64_t around_players = shift<NORTH>(players) |
5          shift<SOUTH>(players) | shift<EAST>(players) |
6          shift<WEST>(players) | shift<NORTH_EAST>(players) |
7          shift<NORTH_WEST>(players) | shift<SOUTH_EAST>(players) |
8          shift<SOUTH_WEST>(players);
9      return (around_players & possible) == 0;
10 }
```

---

Listing 6: Game over test

Now that we know how to test for game over, we will study how to efficiently generate possible moves in a given position.

## 2.3 Generating Possible Piece Moves

Move generation will use a technique called *magic bitboards*. To understand this technique, I will first present magic perfect hashing through a simple example. I drew inspiration from the article [http://pradu.us/old/Nov27\\_2008/Buzz/research/magic/Bitboards.pdf](http://pradu.us/old/Nov27_2008/Buzz/research/magic/Bitboards.pdf) for this example.

### 2.3.1 Magic Perfect Hashing Function

Suppose we need to find the position of the only 1-bit in an unsigned long integer (`uint64_t`). For example, if we call `position` the function that calculates this position, we get the following results:

```
position(1)           == 0
position(0b1000)      == 3
position(0b10000000) == 7
position(0x8000000000000000) == 63
```

Note that processors have instructions to efficiently compute this function<sup>5</sup>, and the gcc compiler allows efficient calculation of this position with the function `__builtin_ctzll(unsigned long long x)`<sup>6</sup>, for example,

```
position(0b1000) == __builtin_ctzll(0b1000)
```

However, we will code this `position` function using a magic perfect hashing function. This function will be useful for move generation later. Using a standard hash table, we could first fill it with 64 values and then simply look up this table, as in the following code:

---

```
1 unordered_map<uint64_t, uint8_t> table {
2     {1, 0}, {0b10, 1}, {0b100, 2}, {0b1000, 3}, //...
3 }
4 int position(uint64_t x) {
5     return table[x];
6 }
```

---

Looking up values in this hash table is far more costly than indexing a simple array<sup>7</sup>. However, we cannot directly use an array because the indexing values span too large a range—for example, the value `0x8000000000000000` (9, 223, 372, 036, 854, 775, 808!). The idea is to find an efficient hash function that transforms each of the 64 values<sup>8</sup>—which we call bitboards—into the range  $[0, 2^6 - 1 = 63]$ . Moreover, we want no collisions—this is called perfect hashing. If there were collisions, positions would be incorrect. For example, if keys `0b100` and `0b1000000` were both transformed to index 42, we would need to store value 2 or 6 in `table[42]`, but we cannot store more than one value at the same location.

The magic perfect hashing function will have the following form:

---

```
1 constexpr uint64_t MAGIC = //...
2 constexpr int K = 6;
3 int magic_perfect_hashing(uint64_t bitboard) {
4     return bitboard * MAGIC >> (64 - K);
5 }
```

---

---

<sup>5</sup>`tzcnt` on Intel.

<sup>6</sup>Returns the number of trailing 0-bits in `x`, starting at the least significant bit position. If `x` is 0, the result is undefined.

<sup>7</sup>We will measure the performance gain from magic perfect hashing compared to using a standard hash table at the end of this section.

<sup>8</sup>`1`, `0b10`, `0b100`, `0b1000`, `0b10000`, ....

In this function:

- On line 2, the constant `K` gives the number of bits for our index. Here with `K == 6`, this gives us a maximum of  $2^K = 2^6 = 64$  possible values in the table.
- On line 4, we see the formula for calculating the table index from the `uint64_t` `key` parameter. We multiply the key by the `MAGIC` constant and then right-shift to keep only the `K` most significant bits of the result. This operation is very inexpensive, but how do we find this magic constant?

A simple method to find the `MAGIC` constant is to generate it randomly until we find a value that produces no collisions! This approach is used in Stockfish [3] to generate magic bitboards. Listing 7 shows this approach.

---

```
1  #include <bits/stdc++.h>
2
3  using namespace std;
4
5  int main() {
6      random_device rd;
7      mt19937_64 mt(rd());
8      uniform_int_distribution<uint64_t> d;
9      unordered_map<uint64_t, int> positions;
10     for (int i = 0; i < 64; i++) {
11         positions[1ULL << i] = i;
12     }
13     constexpr int K = 7;
14     auto index = [](uint64_t magic, int k, uint64_t bitboard) {
15         return bitboard * magic >> (64 - k);
16     };
17     while (true) {
18         uint64_t MAGIC = d(mt);
19         bool found = true;
20         set<uint64_t> seen;
21         for (const auto [bitboard, pos] : positions) {
22             int64_t i = index(MAGIC, K, bitboard);
23             if (seen.contains(i)) {
24                 found = false;
25                 break;
26             }
27             seen.insert(i);
28         }
29         if (found) {
```

```
30         cout << format("found magic for K = {}: {:#x}\\n", K, MAGIC);
31         int size = 1 << K;
32         vector<int> table(size, -1);
33         for (const auto [bitboard, pos] : positions) {
34             table[index(MAGIC, K, bitboard)] = pos;
35         }
36         cout << format("uint8_t positions[{}] = {{", size);
37         for (int i = 0; i < size; i++) {
38             cout << table[i] << ',';
39         }
40         cout << "};\\n";
41         break;
42     }
43 }
44 }
```

---

Listing 7: Random search for the **MAGIC** constant

- On lines 6 to 8, we set up a random generator to randomly generate `uint64_t` values.
- On line 9, the lookup table `positions` associates each long integer containing only one 1-bit with its position. This table is initialized on lines 10 to 12.
- On line 13, the constant `K` is the number of bits in our key obtained by magic perfect hashing. `K` must be large enough to represent all values we need to cover. For example, here we have 64 possible values since we have 64 possible positions for the 1-bit in a 64-bit long integer. With `K == 7`, we can cover  $2^7 = 128$  different values. Note that this is more than the 64 values we need, and `K == 6` would suffice (`K == 5` would be too small). However, when trying to find the **MAGIC** constant randomly for `K == 6`, this program could not find one in reasonable time ☹ (we will see another approach to succeed with `K == 6` ☺).
- On lines 14 to 16, the `index` function calculates the index for key `bitboard` (this `bitboard` contains only one 1-bit) using the magic perfect hashing formula.
- The `while` loop, lines 17 to 43, loops until it finds a **MAGIC** constant that achieves magic perfect hashing (this loop could potentially run forever).
- On line 18, we create the **MAGIC** value randomly.
- The `set<uint64_t>` on line 20 allows us to remember the indices (obtained via the `index` function) we have already used, to verify we have no collisions.
- On lines 21 to 28, the `for` loop tests all `bitboards`, finds the index for each using the hash (`index` function), and verifies there are no collisions (line 23) with indices obtained for previous `bitboards`.



- Finally, on lines 29 to 42, if we found a **MAGIC** constant that creates magic perfect hashing, we display it and then display a *C++* array containing the position of the 1-bit for each bitboard. The output of this program for a given execution is shown below.

---

```
found magic for K = 7: 0x65e4d4ee86638416
uint8_t positions[128] = {
    63, -1, 54, -1, 49, 55, 33, -1, 50, -1, -1, 56, 34, -1, 43, -1,
    51, -1, -1, 11, -1, -1, 57, 3, 39, 35, 14, -1, 44, 22, -1, -1,
    52, 31, -1, -1, -1, -1, 12, 20, -1, 18, -1, -1, 58, -1, -1, 4,
    60, 40, 0, 36, -1, 15, -1, -1, 45, -1, 27, 23, 6, -1, -1, -1,
    62, 53, 48, 32, -1, -1, -1, 42, -1, 10, -1, 2, 38, 13, 21, -1,
    30, -1, -1, 19, 17, -1, -1, -1, 59, -1, -1, -1, -1, 26, 5, -1,
    61, 47, -1, 41, 9, 1, 37, -1, 29, -1, 16, -1, -1, -1, 25, -1,
    46, -1, 8, -1, 28, -1, -1, 24, -1, 7, -1, -1, -1, -1, -1, -1
};
```

---

To get the position of the 1-bit for bitboard `0b1000` for example, we need to look up the following value:

---

```
positions[0b1000 * 0x65e4d4ee86638416 >> (64 - 7)]
== positions[0x2f26a774331c20b0 >> 57]
== positions[0x17]
== positions[23]
== 3
```

---

Notice that there are vacant slots in the `positions` array—those containing `-1`—so we wasted space. We would not have wasted any space if we had found a **MAGIC** constant for `K == 6`. But is this possible?

To answer this question, we will proceed differently. Obviously, we do not want to scan through all  $2^{64}$  (9, 223, 372, 036, 854, 775, 808) possible values and test each one for magic perfect hashing. We will use an Satisfiability Modulo Theories (SMT) solver that allows us to set constraints before searching the entire space. Setting constraints first will prune the search space. Note that we have no guarantee the search will be fast, but for this problem, finding a **MAGIC** constant for `K == 6` will be very quick. If you are interested in how an SMT solver works, you can consult [4] and [5].

The code in Listing 8 describes the approach using the Z3 Theorem Prover (Z3) solver [6].

```
1  from z3 import *
2
3  positions = {}
4  for i in range(64):
5      positions[1 << i] = i
6
7  solver = Solver()
8  MAGIC = BitVec('magic', 64)
9  K = 6
10 bitboards = list(positions.keys())
11
12 def index(magic, k, bitboard):
13     return magic * bitboard >> (64 - k)
14
15 for i in range(64):
16     index1 = index(MAGIC, K, bitboards[i])
17     for j in range(i + 1, 64):
18         index2 = index(MAGIC, K, bitboards[j])
19         solver.add(index1 != index2)
20
21 if solver.check() == sat:
22     model = solver.model()
23     m = model[MAGIC].as_long()
24     print(f'found magic for K = {K}: {m:#x}')
25     size = 1 << K
26     table = [-1] * size
27     for bitboard, pos in positions.items():
28         table[index(m, K, bitboard) & size - 1] = pos
29     print(f'constexpr uint8_t positions[{size}] = {{{')
30     for i in range(size):
31         print(f'{{table[i]}}, ', end='')
32     print('\n};')
```

---

Listing 8: Searching for the MAGIC constant with an SMT solver

- Lines 3 to 5 associate each bitboard containing only one 1-bit with the corresponding position of that bit. We use the `positions` dictionary (line 3) for this.
- On line 7, we instantiate the Z3 solver.
- On line 8, we declare that the `MAGIC` constant is of type `z3.BitVec`, a type provided by the Z3 solver for representing bit vectors—here we use a 64-bit vector.

- The `index` function defined on lines 12 and 13 calculates the hash based on the `MAGIC` constant and `K`.
- On lines 15 to 19, we give the solver the constraints that hashed values from the `index` function for two different bitboards must not map to the same location. This constraint ensures perfect hashing.
- On line 21, we run the solver, which returns `sat` if it found a `MAGIC` value satisfying the constraints. Note that if the solver returns `unsat` instead of `sat`, it means no 64-bit integer allows perfect hashing. Even though the search space is very large, on my machine, finding a solution takes less than a second!
- Lines 22 to 32 print to standard output the `MAGIC` value and the *C++* array of obtained values. This array is shown below.

---

```
found magic for K = 6: 0x2643c51ab9dfa5b
constexpr uint8_t positions[64] = {
    0,  1,  2, 14,  3, 22, 28, 15, 11,  4, 23, 55,  7, 29, 41, 16,
    12, 26, 53,  5, 24, 33, 56, 35, 61,  8, 30, 58, 37, 42, 17, 46,
    63, 13, 21, 27, 10, 54,  6, 40, 25, 52, 32, 34, 60, 57, 36, 45,
    62, 20,  9, 39, 51, 31, 59, 44, 19, 38, 50, 43, 18, 49, 48, 47
};
```

---

Note that in *Python* it was important to perform the operation `index(m, K, bitboard) & size - 1` on line 28 because Python integers have arbitrary precision. For example, we have:

---

```
0x2643c51ab9dfa5b * 0x8000000000000000 >> 58
== 0x4c878a3573bf4b60

(0x2643c51ab9dfa5b * 0x8000000000000000 >> 58) & size - 1
== 0x4c878a3573bf4b60 & 0b111111
== 32
```

---

In *C++* we would have:

---

```
0x2643c51ab9dfa5b * 0x8000000000000000 >> 58
== 32
```

---

Note that we don't have the same problem on line 18 because `MAGIC` is of type `BitVec(64)`.

To test the efficiency of our magic perfect hashing technique, we set up the micro-benchmark shown in Listing 9. We use the *Google Benchmark* library for these measurements [7]. The benchmark execution results are shown in Listing 10. The `BM_magic_positions` function using magic perfect hashing is much faster than the `BM_unordered_map_positions` function using a standard hash table. Of course, the `BM_builtin_ctz_positions` function using a dedicated instruction is the most efficient (but will be of no use for move generation).

---

```
1  #include <benchmark/benchmark.h>
2  #include <unordered_map>
3  #include <vector>
4
5  std::vector<uint64_t> generate_isolated_bits() {
6      std::vector<uint64_t> samples;
7      for (int i = 0; i < 64; i++) {
8          samples.push_back(1ULL << i);
9      }
10     return samples;
11 }
12 static void BM_magic_positions(benchmark::State& state) {
13     constexpr uint64_t MAGIC = 0x2643c51ab9dfa5b;
14     constexpr int K = 6;
15     constexpr uint8_t positions[64] = {
16         0,  1,  2, 14,  3, 22, 28, 15, 11,  4, 23, 55,  7, 29, 41, 16,
17         12, 26, 53,  5, 24, 33, 56, 35, 61,  8, 30, 58, 37, 42, 17, 46,
18         63, 13, 21, 27, 10, 54,  6, 40, 25, 52, 32, 34, 60, 57, 36, 45,
19         62, 20,  9, 39, 51, 31, 59, 44, 19, 38, 50, 43, 18, 49, 48, 47
20     };
21     auto samples = generate_isolated_bits();
22     size_t idx = 0;
23     for (auto _ : state) {
24         uint64_t bitboard = samples[idx++ & 63];
25         benchmark::DoNotOptimize(positions[bitboard * MAGIC >> (64 - K)]);
26     }
27     state.SetItemsProcessed(state.iterations());
28 }
29 BENCHMARK(BM_magic_positions);
30 static void BM_unordered_map_positions(benchmark::State& state) {
31     std::unordered_map<uint64_t, uint8_t> map;
32     for (uint8_t i = 0; i < 64; i++) {
33         map[1ULL << i] = i;
```

```

34     }
35     auto samples = generate_isolated_bits();
36     size_t idx = 0;
37     for (auto _ : state) {
38         uint64_t bitboard = samples[idx++ & 63];
39         benchmark::DoNotOptimize(map[bitboard]);
40     }
41     state.SetItemsProcessed(state.iterations());
42 }
43 BENCHMARK(BM_unordered_map_positions);
44 static void BM_builtin_ctz_positions(benchmark::State& state) {
45     auto samples = generate_isolated_bits();
46     size_t idx = 0;
47     for (auto _ : state) {
48         uint64_t bitboard = samples[idx++ & 63];
49         benchmark::DoNotOptimize(__builtin_ctzll(bitboard));
50     }
51     state.SetItemsProcessed(state.iterations());
52 }
53 BENCHMARK(BM_builtin_ctz_positions);
54 BENCHMARK_MAIN();

```

Listing 9: Micro-benchmark comparing execution times of magic perfect hashing, a standard hash table, and a dedicated machine instruction.

---

1	Run on (12 X 4400 MHz CPU s)				
2	CPU Caches: L1 Data 32 KiB (x6), L1 Instruction 32 KiB (x6),				
3	L2 Unified 256 KiB (x6), L3 Unified 12288 KiB (x1)				
4	Load Average: 0.62, 1.33, 2.92				
5					
6	Benchmark	Time	CPU	Iterations	Throughput
7	-----				
8	BM_magic_positions	0.542 ns	0.542 ns	1000000000	1.85 G/s
9	BM_unordered_map_positions	30.2 ns	30.2 ns	23361655	33.1 M/s
10	BM_builtin_ctz_positions	0.480 ns	0.480 ns	1000000000	2.08 G/s

---

Listing 10: Micro-benchmark execution results from Listing 9

Now that we have seen how the magic perfect hashing technique works through a simple example, we can study generating possible piece moves using this technique.

## 2.3.2 Magic Bitboards for Piece Moves

## 2.3.3 Making and Unmaking Moves

## 2.4 Complete Commented Game Board Code

---

```
1  #ifndef TYPES_H
2  #define TYPES_H
3  #include <stdint>
4  #include <bit>
5  enum Square : int8_t {
6      SQ_A1, SQ_B1, SQ_C1, SQ_D1, SQ_E1, SQ_F1, SQ_G1, SQ_H1,
7      SQ_A2, SQ_B2, SQ_C2, SQ_D2, SQ_E2, SQ_F2, SQ_G2, SQ_H2,
8      SQ_A3, SQ_B3, SQ_C3, SQ_D3, SQ_E3, SQ_F3, SQ_G3, SQ_H3,
9      SQ_A4, SQ_B4, SQ_C4, SQ_D4, SQ_E4, SQ_F4, SQ_G4, SQ_H4,
10     SQ_A5, SQ_B5, SQ_C5, SQ_D5, SQ_E5, SQ_F5, SQ_G5, SQ_H5,
11     SQ_A6, SQ_B6, SQ_C6, SQ_D6, SQ_E6, SQ_F6, SQ_G6, SQ_H6,
12     SQ_A7, SQ_B7, SQ_C7, SQ_D7, SQ_E7, SQ_F7, SQ_G7, SQ_H7,
13     SQ_A8, SQ_B8, SQ_C8, SQ_D8, SQ_E8, SQ_F8, SQ_G8, SQ_H8,
14     SQ_NONE,
15     SQUARE_ZERO = 0,
16     SQUARE_NB   = 64
17 };
18 enum Direction : int8_t {
19     NORTH = 8,
20     EAST  = 1,
21     SOUTH = -NORTH,
22     WEST  = -EAST,
23     NORTH_EAST = NORTH + EAST,
24     SOUTH_EAST = SOUTH + EAST,
25     SOUTH_WEST = SOUTH + WEST,
26     NORTH_WEST = NORTH + WEST
27 };
28 enum File : uint8_t {
29     FILE_A, FILE_B, FILE_C, FILE_D, FILE_E, FILE_F, FILE_G, FILE_H,
30     FILE_NB
31 };
32 enum Rank : uint8_t {
33     RANK_1, RANK_2, RANK_3, RANK_4, RANK_5, RANK_6, RANK_7, RANK_8,
34     RANK_NB
35 };
36 constexpr uint64_t FileABB = 0x0101010101010101;
```

```
37 constexpr uint64_t FileBBB = FileABB << 1;
38 constexpr uint64_t FileCBB = FileABB << 2;
39 constexpr uint64_t FileDBB = FileABB << 3;
40 constexpr uint64_t FileEBB = FileABB << 4;
41 constexpr uint64_t FileFBB = FileABB << 5;
42 constexpr uint64_t FileGBB = FileABB << 6;
43 constexpr uint64_t FileHBB = FileABB << 7;
44 constexpr uint64_t Rank1BB = 0xFF;
45 constexpr uint64_t Rank2BB = Rank1BB << (8 * 1);
46 constexpr uint64_t Rank3BB = Rank1BB << (8 * 2);
47 constexpr uint64_t Rank4BB = Rank1BB << (8 * 3);
48 constexpr uint64_t Rank5BB = Rank1BB << (8 * 4);
49 constexpr uint64_t Rank6BB = Rank1BB << (8 * 5);
50 constexpr uint64_t Rank7BB = Rank1BB << (8 * 6);
51 constexpr uint64_t Rank8BB = Rank1BB << (8 * 7);
52 constexpr uint64_t BLACK_INITIAL_POSITION =
53 0b10000000'00000000'00000000'00001000'00010000'00000000'00000000'00000001;
54 constexpr uint64_t WHITE_INITIAL_POSITION =
55 0b00000001'00000000'00000000'00010000'00001000'00000000'00000000'10000000;
56 constexpr uint64_t FULL = 0xFFFFFFFFFFFFFFFF;
57 constexpr Square make_square(File f, Rank r) {
58     return Square((r << 3) + f);
59 }
60 inline Square lsb(uint64_t b) {
61     return Square(std::countr_zero(b));
62 }
63 inline Square pop_lsb(uint64_t& b) {
64     const Square s = lsb(b);
65     b &= b - 1;
66     return s;
67 }
68 constexpr bool more_than_one(uint64_t b) {
69     return b & (b - 1);
70 }
71 template<Direction D>
72 constexpr uint64_t shift(uint64_t b) {
73     if constexpr (D == NORTH) return b << 8;
74     else if constexpr (D == SOUTH) return b >> 8;
75     else if constexpr (D == EAST) return (b & ~FileHBB) << 1;
76     else if constexpr (D == WEST) return (b & ~FileABB) >> 1;
77     else if constexpr (D == NORTH_EAST) return (b & ~FileHBB) << 9;
78     else if constexpr (D == NORTH_WEST) return (b & ~FileABB) << 7;
```

```
79     else if constexpr (D == SOUTH_EAST) return (b & ~FileHBB) >> 7;
80     else if constexpr (D == SOUTH_WEST) return (b & ~FileABB) >> 9;
81     else return 0;
82 }
83 constexpr uint64_t shift_all_directions(uint64_t b) {
84     uint64_t b1 = b & ~FileHBB;
85     uint64_t b2 = b & ~FileABB;
86     return b << 8 | b >> 8 | b1 << 1 | b1 << 9
87         | b1 >> 7 | b2 >> 1 | b2 << 7 | b2 >> 9;
88 }
89 #endif
```

---

Listing 11: Squares, files, ranks and directions

---

```
1  #ifndef MOVE_H
2  #define MOVE_H
3  #include "types.h"
4  #include <iostream>
5  class Move {
6      uint16_t data;
7  public:
8      Move() = default;
9      constexpr explicit
10     Move(uint16_t d) : data(d) {}
11     constexpr explicit
12     Move(Square from, Square to) : data((to << 6) + from) {}
13     constexpr Square to_sq() const {
14         return Square((data >> 6) & 0x3F);
15     }
16     constexpr Square from_sq() const {
17         return Square(data & 0x3F);
18     }
19     static constexpr Move none() {
20         return Move(0);
21     }
22     constexpr bool operator==(const Move& m) const {
23         return data == m.data;
24     }
25     constexpr bool operator!=(const Move& m) const {
26         return !(*this == m);
```



```
27     }
28     constexpr explicit operator bool() const {
29         return data != 0;
30     }
31     constexpr uint16_t raw() const {
32         return data;
33     }
34 };
35 std::ostream& operator<<(std::ostream& os, const Move& m);
36 std::istream& operator>>(std::istream& is, Move& m);
37 #endif
```

---

Listing 12: Move implementation

---

```
1  #ifndef GAME_H
2  #define GAME_H
3  #include <iostream>
4  #include <vector>
5  #include <utility>
6  #include "types.h"
7  #include "move.h"
8  #include "json.hpp"
9  #include "misc.h"
10 using json = nlohmann::json;
11 class Yolah {
12     uint64_t black = BLACK_INITIAL_POSITION;
13     uint64_t white = WHITE_INITIAL_POSITION;
14     uint64_t empty = 0;
15     uint8_t black_score = 0;
16     uint8_t white_score = 0;
17     uint8_t ply = 0;
18 public:
19     static constexpr uint8_t MAX_NB_MOVES = 75;
20     static constexpr uint8_t MAX_NB_PLIES = 120;
21     class MoveList {
22         Move moveList[MAX_NB_MOVES], *last;
23     public:
24         explicit MoveList() : last(moveList) {}
25         const Move* begin() const { return moveList; }
26         const Move* end() const { return last; }
```

```
27     Move* begin() { return moveList; }
28     Move* end() { return last; }
29     size_t size() const { return last - moveList; }
30     const Move& operator[](size_t i) const { return moveList[i]; }
31     Move& operator[](size_t i) { return moveList[i]; }
32     Move* data() { return moveList; }
33     friend class Yolah;
34 };
35 static constexpr uint8_t BLACK = 0;
36 static constexpr uint8_t WHITE = 1;
37 static constexpr uint8_t HOLE = 2;
38 static constexpr uint8_t FREE = 3;
39 constexpr std::pair<uint8_t, uint8_t> score() const;
40 constexpr int8_t score(uint8_t player) const;
41 constexpr uint8_t current_player() const;
42 static constexpr uint8_t other_player(uint8_t player);
43 uint8_t get(Square) const;
44 uint8_t get(File f, Rank r) const;
45 bool game_over() const;
46 void play(Move m);
47 void undo(Move m);
48 void moves(uint8_t player, MoveList& moves) const;
49 void moves(MoveList& moves) const;
50 bool valid(Move m) const;
51 constexpr uint64_t free_squares() const;
52 constexpr uint64_t occupied_squares() const;
53 constexpr uint64_t bitboard(uint8_t player) const;
54 constexpr uint16_t nb_plies() const;
55 constexpr uint64_t empty_bitboard() const;
56 std::string to_json() const;
57 static Yolah from_json(std::istream& is);
58 static Yolah from_json(const std::string&);
59 friend std::ostream& operator<<(std::ostream& os, const Yolah&
↪ yolah);
60 };
61 std::ostream& operator<<(std::ostream& os, const Yolah& yolah);
62 bool Yolah::game_over() const {
63     uint64_t possible = ~empty & ~black & ~white;
64     uint64_t players = black | white;
65     uint64_t around_players = shift<NORTH>(players) |
↪ shift<SOUTH>(players) | shift<EAST>(players) |
```

```
66         shift<WEST>(players) | shift<NORTH_EAST>(players) |  
↪ shift<NORTH_WEST>(players) |  
67         shift<SOUTH_EAST>(players) | shift<SOUTH_WEST>(players);  
68     return (around_players & possible) == 0;  
69 }  
70 #endif
```

---

Listing 13: Board implementation

## 2.5 Testing with Random Games

## 2.6 What's Next



# Chapter 3

## AI Players



## Chapter 4

### Monte Carlo Player





## Chapter 5

### MCTS Player



## Chapter 6

### Minmax Player



## Chapter 7

# Minmax with Neural Network Player



## Chapter 8

# AI Tournament





## Chapter 9

## Conclusion



# Acronymes

**SMT** Satisfiability Modulo Theories. 23, 24

**Z3** Z3 Theorem Prover. 23, 24



# Bibliography

- [1] Anthropic. *Claude Sonnet 4.5*. <https://www.anthropic.com/claude>. Large Language Model. 2025.
- [2] cppreference.com. *C++ Reference*. Accessed: 2025-01-28. 2025. URL: <https://en.cppreference.com/>.
- [3] Stockfish Team. *Stockfish: Open Source Chess Engine*. <https://stockfishchess.org/>. Accessed: 2025-01-28. 2025.
- [4] Daniel Kroening and Ofer Strichman. *Decision Procedures: An Algorithmic Point of View*. 2nd. Texts in Theoretical Computer Science. An EATCS Series. Berlin Heidelberg: Springer-Verlag, 2016. ISBN: 978-3-662-50496-3.
- [5] Armin Biere, Marijn Heule, Hans van Maaren, and Toby Walsh, eds. *Handbook of Satisfiability*. 2nd. Vol. 336. Frontiers in Artificial Intelligence and Applications. IOS Press, 2021. ISBN: 978-1-64368-160-3.
- [6] Leonardo de Moura and Nikolaj Bjørner. “Z3: An Efficient SMT Solver”. In: *Tools and Algorithms for the Construction and Analysis of Systems (TACAS)*. Vol. 4963. Lecture Notes in Computer Science. Berlin, Heidelberg: Springer, 2008, pp. 337–340.
- [7] Google. *Google Benchmark: A microbenchmark support library*. <https://github.com/google/benchmark>. C++ microbenchmarking library. 2025.